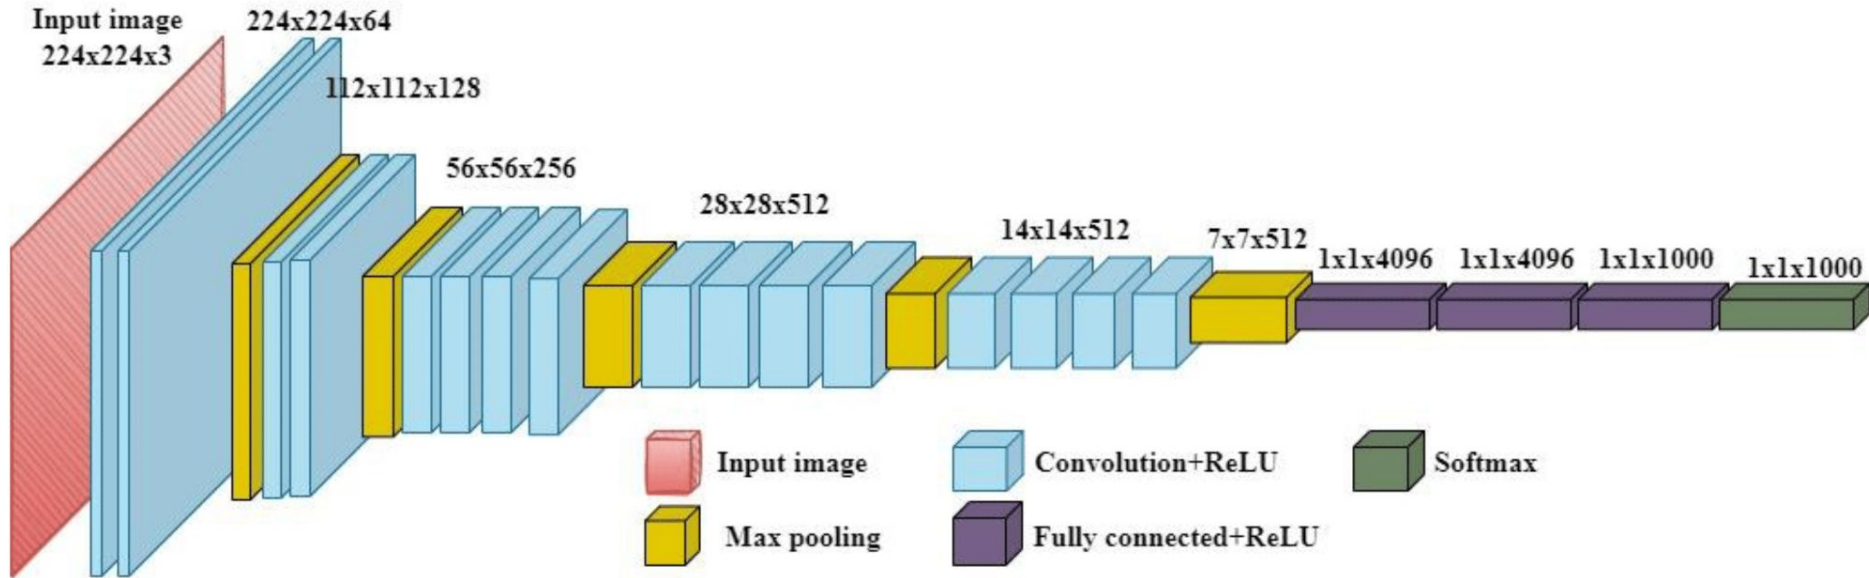





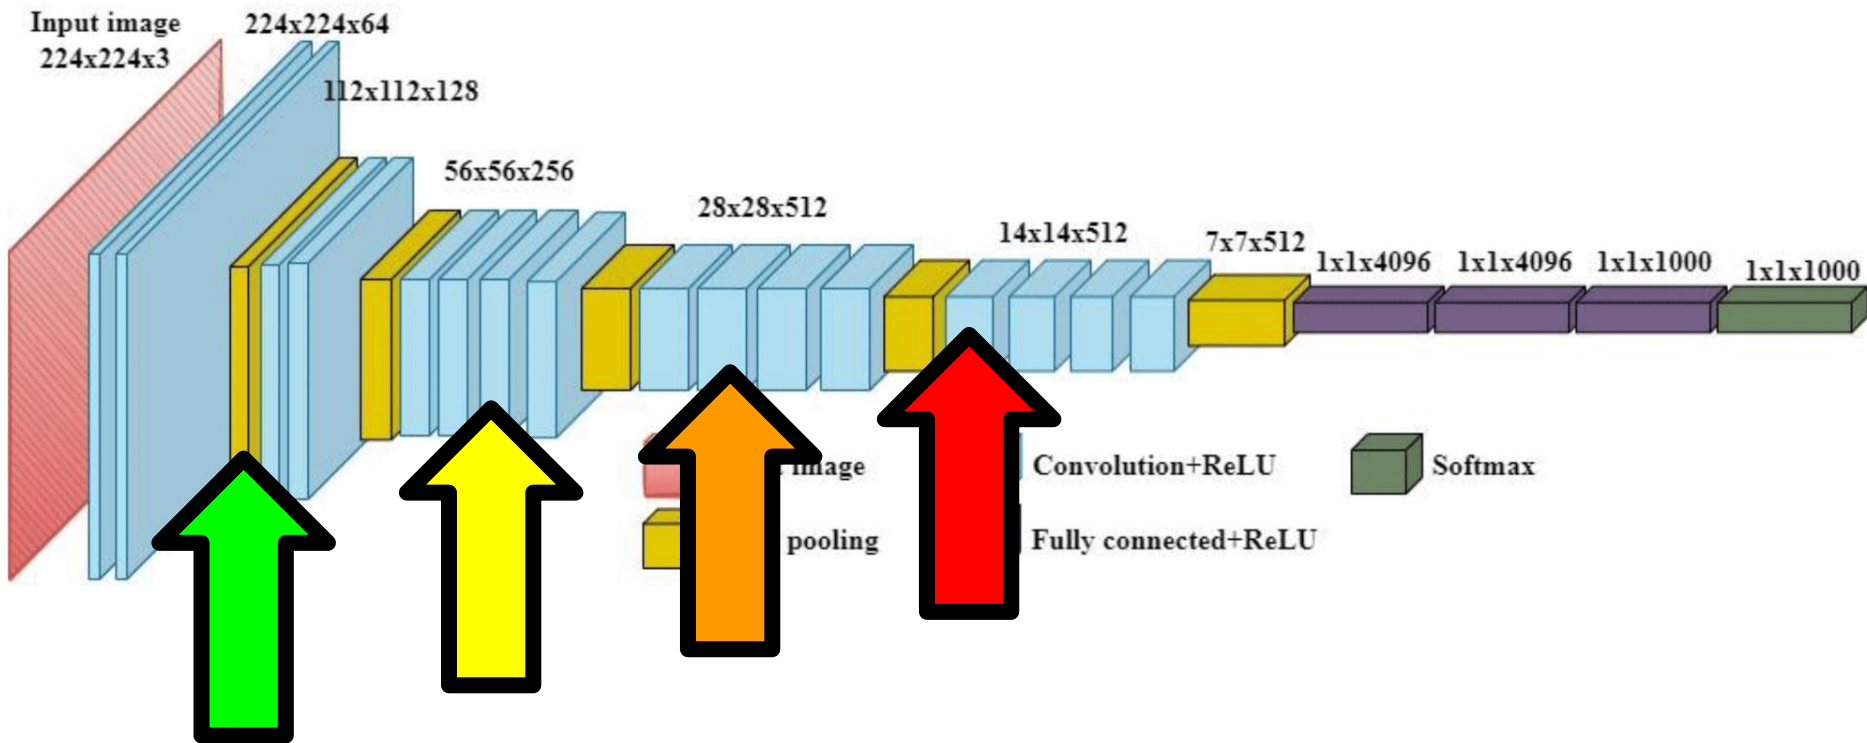


Image classifier






compression



layer=2



layer=5



layer=8



layer=11



layer=14




layer=17



layer=20




layer=23




layer=26




layer=29





layer=32




layer=35









Gram matrix (aka. covariance aka. correlation)

The diagram illustrates the calculation of a Gram matrix. On the left, four vertical vectors (yellow, green, magenta, black) are multiplied by their corresponding four horizontal vectors (yellow, green, magenta, black) using the dot product operator '@'. This results in a 4x4 matrix of values, where each row and column is color-coded to match the vectors it represents.

	Yellow	Green	Magenta	Black
Yellow	1	0.2	0.7	0.4
Green	0.2	1	0.1	0.3
Magenta	0.7	0.1	1	0.9
Black	0.4	0.3	0.9	1



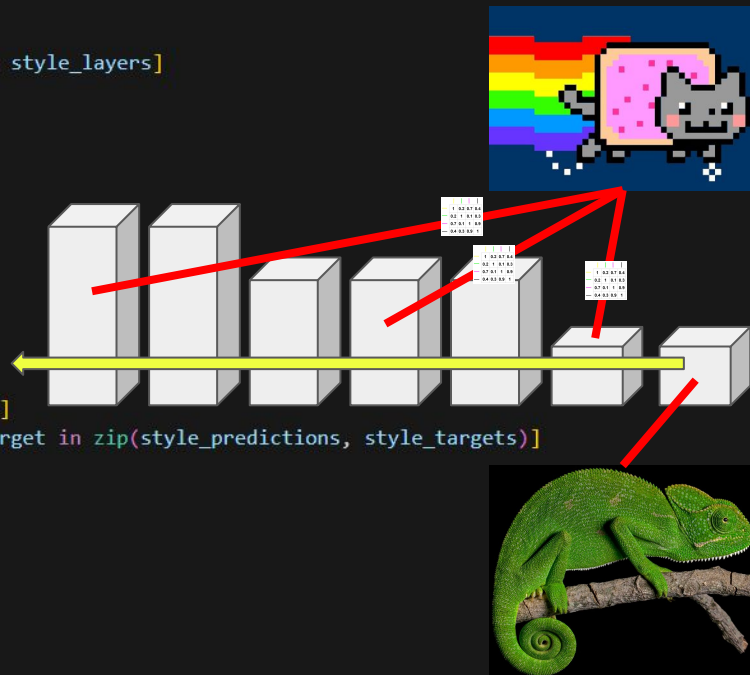




```
style_layers = [1, 6, 11, 20, 29]
content_layer = 28
```

```
def style_transfer(start, content, style, scaler=1., content_layer=content_layer, style_layers=style_layers, epochs=10, m=vgg_hooked, mem=vgg_mem):
    content_activations = save_activations(content)
    style_activations = save_activations(style)
    # move to device
    start = copy.deepcopy(start.detach()).to(device).requires_grad_()
    content_target = content_activations[content_layer].to(device)
    style_targets = [gram(style_activations[layer]).to(device) for layer in style_layers]
    m = m.to(device)
    m.eval()
    # optimizer
    optimizer = torch.optim.LBFGS([start])
    for epoch in tqdm(range(epochs)):
        def closure():
            m(start)
            # content loss
            content_predicted = mem[content_layer]
            content_loss = F.mse_loss(content_predicted, content_target)
            # style loss
            style_predictions = [gram(mem[layer]) for layer in style_layers]
            style_losses = [F.mse_loss(predicted, target) for predicted, target in zip(style_predictions, style_targets)]
            style_loss = torch.stack(style_losses).sum()
            # merge losses
            loss = content_loss + scaler * style_loss
            optimizer.zero_grad()
            loss.backward()
            return loss

        optimizer.step(closure)
        start.data = torch.clip(start, 0.0, 1.0).data
        if epoch % log_every == log_every - 1:
            print(epoch)
    return start
```



result



content



style



result



content



style



result



content



style



swe-to-mle.pages.dev